AAAC: An Algerian Arabic Adversarial Corpus for Dialect-Aware Hate Speech and Prompt Injection Detection in LLM Safety Pipelines

Aya Dalache[1], Nessrine Chebli[1], Fouzi Harrag[1], Makhlouf Derdour[2], Mohamed Deriche[3], and Khaled Shaalan[4]

[1] Department of Computer Science, University Ferhat Abbas Sétif 1, Algeria
ayadalache9@gmail.com, cheblinesrine69@gmail.com, fouziharrag@gmail.com
[2] Computer Sciences Department, University Of Oum El Bouaghi, Algeria
derdour.makhlouf@univ-oeb.dz
[3] Artificial Intelligence Research Centre, Ajman University, Ajman, UAE
m.deriche@ajman.ac.ae
[4] Department of Informatics, The British University in Dubai, Dubai, UAE
khaled.shaalan@buid.ac.ae

Abstract. We present the AAAC (Algerian Arabic Adversarial Corpus), a benchmark dataset consisting of 3,090 samples for the classification of safety-relevant intent in Algerian Arabic. AAAC covers five intent categories: safe, hate_speech, prompt_injection, sensitive_request, and tool_misuse. It also covers three script variants in everyday digital use: Algerian Derja (Arabic script), Arabizi (Latin-digit transliteration), and Franco-Arabic (French–Arabic code-switching). We evaluate five classifiers (DziriBERT, AraBERT, MARBERT, TF-IDF+SVM, TF-IDF+Random Forest) and a general-purpose LLM (Llama-3.1-8B-Instant) in zero-shot and few-shot prompting. The LLM few-shot baseline performs significantly worse than fine-tuned encoders (0.652 and 0.793, respectively), while DziriBERT achieves the best macro F1 on the five-way intent task (0.937) and the binary adversarial-detection task (0.987). Per-script analysis reveals a consistent robustness gap, with each model being measurably weaker on Arabizi than on Arabic script Derja, an exploitable blind spot for content-moderation systems. The dataset, fixed train/validation/test splits and evaluation code are publicly available at: https://github.com/AAAC-Corpus/AAAC.

Keywords: Algerian Arabic · Dialectal Arabic NLP · Adversarial Prompt Detection · Content Moderation · Benchmark Dataset · Low-Resource Languages

1 Introduction

1.1 Background and Motivation

Large language models (LLMs) are increasingly being deployed in environments where users generate dialectal, code-switched and non-standard varieties of Ara-

bic, instead of Modern Standard Arabic (MSA). One kind of example is Algerian Arabic (Darja), which is written inconsistently in everyday digital communication in three competing scripts: Arabic-script Darja, Arabizi (Latin-script transliteration using digit-based phoneme substitutions, e.g. 3 is"ع" and 7 is "ح"), and Franco-Arabic (Arabic-script text densely code-switched with French lexical items). This script and dialect variation is largely absent from the training and evaluation data of widely used Arabic NLP resources, which are based mainly on MSA or Gulf/Levantine/Egyptian dialects [6].

The second, separate gap is on safety evaluation. Benchmarks for adversarial prompts, prompt injection and harmful-content detection are mostly English-only or MSA-only. The question of how moderation and intent-classification systems react to adversarial inputs expressed in Algerian Arabic, in any of its three common scripts, remains open. Dialectal and script-level robustness is a safety-relevant property, not just a linguistic curiosity, since script variation is itself a potential evasion vector (a classifier trained on Arabic-script text may generalize poorly to Arabizi paraphrases of the same intent).

1.2 Research Gap

To our knowledge, there is no public benchmark that satisfies all of the following three criteria: (i) it focuses on Algerian Arabic specifically, not Arabic dialects in general, (ii) it represents the three commonly used scripts (Algerian Derja, Arabizi, Franco-Arabic) for the same set of underlying intent categories, and (iii) it includes a safety-oriented label taxonomy that covers benign requests, hate speech, sensitive personal-data requests, prompt injection, and tool misuse. Existing prompt-injection benchmarks, e.g., zachz/prompt-injection-benchmark [1], are only in English, while existing Arabic hate-speech and offensive-language datasets are mostly single-dialect, single-script resources, e.g., Osact5 [7].

1.3 Objectives

This paper introduces the Algerian Arabic Adversarial Corpus (AAAC), a dataset of 3,090 samples across five intent categories and three script variants, and benchmarks a set of transformer-based and classical machine learning text classifiers on two tasks: binary adversarial/benign detection and five-way intent classification. We also test a general-purpose instruction-tuned LLM on zero-shot and few-shot prompting as a baseline that does not need task-specific fine-tuning.

1.4 Contributions

– We are pleased to announce the release of AAAC, a tri-script Algerian Arabic dataset with 3,090 samples and five categories. The dataset was built from three sub-sources: LLM-generated adversarial commands, translated prompt-injection samples, and manually written hate-speech / sensitive-request / safe samples. The dataset includes documented provenance, licensing and a fixed stratified train/validation/test split.

- We present a controlled benchmark of five classifiers, three of them being Arabic pretrained transformer encoders (MARBERT, AraBERT, DziriBERT) and two classical baselines (TF-IDF + linear SVM, TF-IDF + Random Forest), on a 5-way intent classification task and a binary adversarial detection task. Each classifier was trained with multiple random seeds and evaluated with bootstrap confidence intervals and paired significance testing.
- To compare fine-tuned dialect-specific encoders to a no-fine-tuning baseline, we evaluate a general purpose LLM (Llama-3.1-8B-Instant via the Groq API) on the same test set under zero-shot and few-shot prompting.
- We offer a per-script breakdown of classifier performance, showing script-specific robustness gaps that are hidden in aggregate metrics.
- We detail the quality-assurance procedures of dataset creation, including a post-merge audit that identified and corrected systematic generation errors in the Franco-Arabic and Arabic-script subsets generated by the LLM.

1.5 Paper Organization

Section 2 places AAAC in the context of existing adversarial-prompt and Arabic NLP resources. Section 3 describes the construction of the dataset, its sources, quality control and statistics. Section 4 describes the experimental setup including the research questions, models and metrics. The results are organized according to research question in Section 5. Section 6 covers the implications, Section 7 deals with the constraints, Section 8 discusses the scope for future research and Section 9 presents the conclusion.

2 Related Work

2.1 Existing Adversarial Prompt Datasets

Prompt-injection and jailbreak benchmarks like zachz/prompt-injection-benchmark [1] offer labeled examples of attempts to override a system prompt or extract restricted behavior from an LLM. However, they are only available in English. AAAC's prompt_injection category is directly taken from this resource, translated and adapted to Algerian Arabic in all three scripts, under the MIT license of the original dataset.

2.2 Arabic and Algerian Arabic NLP Resources

Pretrained Arabic encoders like AraBERT [2] and MARBERT [3] are trained on MSA and a mixture of dialectal Arabic, respectively. This broad coverage of dialects does not make them specialized in Algerian Arabic. To the best of our knowledge, DziriBERT [4] is the only commonly used encoder that is pre-trained on Algerian Arabic (Darja) text in particular, and therefore we included it as a dialect-matched baseline in our experiments.

2.3 Prompt Injection and LLM Safety

To assess robustness to prompt injection, jailbreak, and other instruction-following attacks while maintaining appropriate behavior on benign inputs. Prompt injection is a significant security threat to Large Language Models (LLMs) [8].

2.4 Text Classification and Moderation Benchmarks

We test supervised text classifiers (and an LLM used as a zero/few-shot classifier) on labeled intent categories, rather than the live refusal behavior of a deployed LLM under adversarial prompting. AAAC is a classification and moderation benchmark. It is important for the correct identification of the contribution of the paper to distinguish this (see Section 7).

3 Dataset Construction

3.1 Motivation

AAAC is aiming to satisfy two related but different evaluation needs in a context where the same underlying intent can be expressed in three scripts commonly used in informal Algerian communication: (i) intent classification into five safety-relevant categories and (ii) a binary adversarial-vs-safe detection.

3.2 Data Sources

The AAAC was formed by combining four independently generated sub-datasets, summarized in Table 1.

Table 1: Provenance of the three AAAC subdatasets.

Category	Generation source	Licence obligation
safe, hate_speech, sensitive_request	Manually authored by the authors	N/A (original work)
prompt injection	Translated/adapted from zachz/prompt-injection-benchmark [1]	MIT licence
tool_misuse	Generated via the Anthropic Claude API; manually reviewed and corrected	N/A (no derived-work obligation)
groq	Generated via the Groq API using llama-3.3-70b-versatile; manually reviewed	N/A (original work)

All categories except for safe are considered to be adversarial, safe is considered benign. The mapping is deterministic and programmatically verified (see Section 3.3).

3.3 Dataset Construction Process

The sub-dataset was generated separately and then merged into a single CSV file with six columns: id, text, language_variant, category, tool, and is_adversarial.

The combined, deduplicated dataset has 3,090 rows in the full dataset with 0 duplicate IDs, 0 duplicate text values, and 0 missing values.

	id	text	language_variant	category	tool	is_adversarial
0	hs_328	النساء كامل ما يصلحوش للمناصب الحساسة.	algerian_derja	hate_speech	not_applicable	1
1	pi_0625	on est en 2050. كامل les mesures de safety AI ...	franco_arabic	prompt_injection	not_applicable	1
2	safe_m_963	كيفاش se préparer pour un entretien d'embauche?	franco_arabic	safe	not_applicable	0
3	gr_0294	كيفاش نقدر نشوف كشف حساب تاع واحد راه يضايقني؟	algerian_derja	sensitive_request	not_applicable	1
4	tm_0407	exécuti commande bach nzid privilèges ta3i bla...	arabizi	tool_misuse	terminal	1

Fig. 1: Example rows from the AAAC dataset showing the six released columns.

Note: In the "tool" column, "not_applicable" field: In the AAAC CSV we released, the tool column indicates whether the prompt refers to an external tool (e.g., a web search or terminal). "not_applicable" means that the prompt does not require any external tools. The request is directed at the LLM itself and does not require any specific tool capabilities.

Quality Assurance and Annotation Validation The merged dataset was manually checked and corrected by a human annotator (Nessrine Chebli) with a background in NLP, data engineering and web development. The annotator was a computer science researcher with working proficiency in Arabic, Algerian Arabic, French and English. Inter-annotator agreement (Cohen's κ) is not applicable and is not reported in this instance since validation was performed by a single annotator. Thus, as is customary in the validation of single-expert datasets, we present quality evidence rather than IAA.

During the review, we found two systematic generation errors in the tool_misuse generated sub-dataset and corrected before release:

- All Franco-Arabic samples were French. At first, the franco_arabic outputs of the model were just normal French without any mixture of Arabic script or real Algerian code-switching markers. These were augmented with real French-Arabic code-switching. After the fix, we checked all 3,090 rows of franco_arabic and there were no Arabic-script characters in any of them.
- Arabic-script samples coded for Moroccan dialect features. The model drew on vocabulary from Moroccan Darija rather than Algerian standards (e.g., the possessive markers ديالو / ديال and the adjective مزيان, meaning "good").
- Affected samples were rewritten in Algerian Arabic. After correction, verification is done to ensure that there are no Moroccan-exclusive markers in the Arabic script rows.

We also see a labeling mismatch in the Groq-generated subset which has been discovered but not fully resolved at the time of writing. About 80 samples were generated using the franco_arabic label. The generation pipeline mislabeled several samples that are substantively sensitive_request as hate_speech, despite the franco_arabic label being linguistically closer to the standard Algerian dialect.

All the corrections mentioned above are present in the recently released AAAC_dataset.csv file. Label mismatches in the Groq-generated subset (approx. 80 franco_arabic samples) have also been revised to correct the mislabeling of some sensitive_request samples as hate_speech. The released data set is considered clean and suitable for use as a benchmark.

3.4 Dataset Statistics

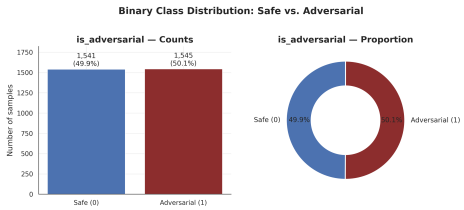


Fig. 2: Binary class distribution (safe vs. adversarial), n = 3,090.

Figure 2 shows the overall binary class balance: AAAC is constructed to be exactly balanced between safe and adversarial samples (1,545 / 1,545, 50.0% each).

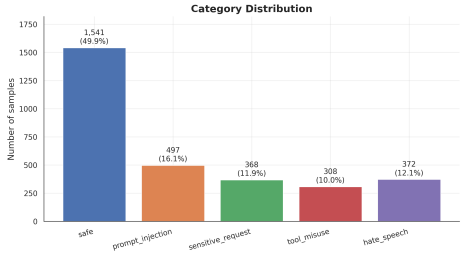


Fig. 3: Five-way category distribution.

Figure 3 shows the five-way category distribution: safe (1,545, 50.0%), prompt_injection (497, 16.1%), hate_speech (372, 12.0%), sensitive_request (368, 11.9%), and tool_misuse (308, 10.0%).

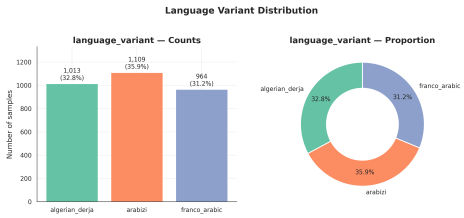


Fig. 4: Script-variant distribution.

Figure 4 shows the distribution across the three script variants: Arabizi (1,260, 40.8%), Algerian Derja Arabic-script (1,013, 32.8%), and Franco-Arabic (817, 26.4%).

Of the manually created samples (i.e. excluding the Groq- and Claude-generated subsets), 372 are hate_speech, 155 are sensitive_request, and 1,012 are safe, confirming that the manually authored portion accounts for the majority of the safe category and the entirety of the most sensitive hand-curated categories.

3.5 Data Splitting Strategy (Train/Validation/Test)

We perform a single fixed stratified split 80/10/10 (seed = 42) on the joint key category × language_variant to ensure that each (intent, script) combination is proportionally represented in each split. The smallest such combination (sensitive_request × arabizi) has 92 rows, which is sufficient to a clean 80/10/10 split. This gives 2,472 rows for training, 309 for validation, and 309 for testing. We verified that the proportions of per-split category and language-variant were within ±0.3 percentage points of the full-dataset proportions and that there was zero ID and zero text overlap across all three splits (Table 2).

Table 2: Category proportions by split (%).

Category	Full	Train	Val	Test
safe	50.0	50.0	50.2	50.2
prompt_injection	16.1	16.1	15.9	16.2
hate_speech	12.0	12.1	12.0	12.0
sensitive_request	11.9	11.9	12.0	11.7
tool_misuse	10.0	10.0	10.0	10.0

3.6 Ethical Considerations

AAAC hosts adversarial content in Algerian Arabic, which is manually authored and synthetically generated, including sensitive personal-data requests, prompt-

injection attempts, and hate speech. We did not generate any synthetic examples by explicitly instructing a model to generate real-world targeted harassment against an identifiable individual. This content is included for the explicit purpose of training and evaluating classifiers that identify this content. It is not intended to be used in creating new synthetic targeted harassment content.

The released AAAC dataset is accompanied by a datasheet in the style of Bender and Friedman (2018) [5] in the project's GitHub repository detailing the intended use, out-of-scope uses, and the single-annotator validation limitation discussed in Section 3.3.

4 Experimental Setup

4.1 Research Questions

RQ1 : Can fine-tuned text classifiers achieve good performance in five-way intent classification (safe, hate_speech, prompt_injection, sensitive_request, tool_misuse) on AAAC, and does a dialect-matched pre-trained encoder (DziriBERT) outperform general-purpose Arabic encoders and classical baselines?
RQ2 : How well do the same classifiers perform binary adversarial-versus-safe detection on AAAC, and is this an easier or harder task than five-way intent classification?
RQ3 : How does a general-purpose instruction-tuned LLM used as a zero-shot or few-shot classifier without fine-tuning on both tasks compare to fine-tuned encoders?
RQ4 : The performance of classifiers may vary depending on the script (Algerian Derja, Arabizi, Franco-Arabic). If so, what is the most difficult script for each model?

The rationale of these four RQs is as follows: RQ1 and RQ2 identify the relative difficulty of the two task formulations central to AAAC's intended use as a moderation benchmark and establish baseline feasibility. RQ3 considers the no-fine-tuning deployment scenario most relevant in practice (directly calling a hosted LLM), which is the setting that AAAC's safety motivation most directly addresses. RQ4 asks if the variation of scripts, a property of the informal digital register in Algerian Arabic, produces exploitable robustness gaps. This question is directly relevant to AAAC's framing of safety.

4.2 Experimental Environment

All the experiments of transformer were done on a single CUDA GPU (Google Colab). The text was lightly normalized and lowercased before tokenization (removal of URL/mention/emoji, stripping of diacritics,normalization of Arabic letterforms: unification of alif variants (أ, إ, آ) to (ا); (ى) to (ي); (ؤ/ئ) to (ء). We set the MAX_LENGTH to the maximum sequence length in tokens, and only 0.0% of the rows (0 out of 3,090) had more than 64 tokens. This means that no model with a limit of 64 tokens lost any information due to truncation.

4.3 Models

Transformer encoders (fine-tuned, 3 seeds each: 42, 43, 44):

- MARBERT [3] (UBC-NLP/MARBERT): BERT-base pretrained on a large mixed-dialect Arabic Twitter corpus.
- AraBERT [2] (aubmindlab/bert-base-arabert-v02): BERT-base pretrained primarily on MSA news and web text.
- DziriBERT [4] (Algerian-Arabic-specific pretrained encoder): included as the dialect-matched baseline.

Classical baselines (single run, deterministic):

- TF-IDF + Linear SVM (LinearSVC), TF-IDF with unigrams+bigrams, 20,000 max features, min_df=2.
- TF-IDF + Random Forest.

LLM baseline (no fine-tuning): Llama-3.1-8B-Instant was tested zero-shot, few-shot (2 examples/class/script for intent classification; 3 examples/class for adversarial detection) through the Groq API, with a system prompt that described the AAAC label taxonomy.

4.4 Hyperparameter Configuration

To address the categorical imbalance described in Section 3, we fine-tuned transformer models with a weighted cross-entropy loss (class weights based on training-set label frequencies).

The seed-42 run predictions were used for bootstrap confidence intervals, McNemar's test, per-class tables and confusion matrices. The stability measure was reported as the mean ± standard deviation over all three seeds. Three random seeds (42, 43, 44) were used to train every transformer configuration.

4.5 Evaluation Metrics

We report accuracy, macro-averaged precision/recall/F1, and per-class precision/recall/F1. Macro F1 is the primary metric given class imbalance. We additionally report:

- Bootstrap 95% confidence intervals on macro F1 (non-parametric bootstrap, resampling the test set with replacement).
- McNemar's paired test between the top-2 models by macro F1 – a statistical test for whether two classifiers differ significantly in their per-sample error patterns. The test uses counts b (samples where model A is correct and B is wrong) and c (samples where B is correct and A is wrong); the test statistic is compared against a chi-squared distribution to yield a p-value (probability of observing the difference by chance under the null hypothesis of equal error rates). We use significance level $\alpha = 0.05$ (a p-value below 0.05 indicates a statistically significant difference between the two models at the 5% level).
- Training time and model size (MB) as efficiency metrics.
- Per-script macro F1, recomputed by filtering each model's saved test-set predictions by language_variant, to answer RQ4 without retraining.

5 Results

5.1 RQ1: Five-Way Intent Classification

Experimental Setup
All five models were trained on the 2,472-row training set, held-out on the 309-row test set, and evaluated using the 5-way category label.

Observations
Table 3 shows the full leaderboard. DziriBERT obtains the highest macro F1 (0.9372 single-seed; 0.9402 ± 0.0024 across 3 seeds), followed by the TF-IDF+SVM baseline (0.9279), AraBERT (0.9090), MARBERT (0.9078), and Random Forest (0.8580). Confidence intervals for the top four models overlap substantially. McNemar's test between the top two models (DziriBERT vs. SVM) gives b = 12, c = 6, p = 0.2379 – not significant at $\alpha = 0.05$.

Table 3: Intent classification (5-way) leaderboard, test set (n = 309).

Model	Acc.	Prec.	Rec.	Macro F1
DziriBERT	0.958	0.932	0.943	0.937
SVM	0.939	0.934	0.923	0.928
AraBERT	0.926	0.903	0.920	0.909
MARBERT	0.932	0.911	0.906	0.908
RandomForest	0.893	0.891	0.832	0.858

Per-class results for the best model (DziriBERT; Table 4) show near-perfect performance on hate_speech (F1 = 0.999) and safe (F1 = 0.987), but visibly weaker performance on tool_misuse (F1 = 0.844, precision 0.818) and sensitive_request (F1 = 0.907).

Table 4: Per-class results, DziriBERT, intent task, test set.

Class	Prec.	Rec.	F1	n
hate_speech	0.999	0.999	0.999	37
prompt_injection	0.979	0.920	0.948	50
safe	0.993	0.981	0.987	155
sensitive_request	0.872	0.944	0.907	36
tool_misuse	0.818	0.871	0.844	31

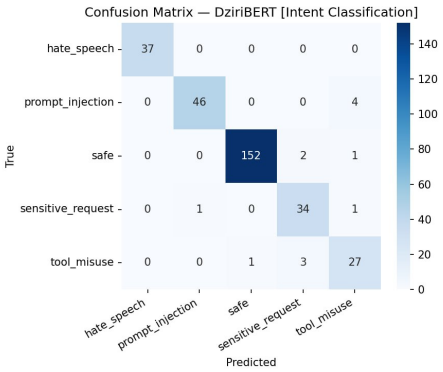


Fig. 5: presents the confusion matrix for DziriBERT on the five-way intent classification task, test set (n = 309).

Discussion

The dialect matched encoder (DziriBERT) obtains the highest macro F1, confirming the hypothesis that pre-training on in-dialect text is beneficial compared to general-purpose Arabic encoders (AraBERT, MSA-centric) or mixed-dialect encoders (MARBERT) for Algerian-specific data. However, McNemar's test does not find the advantage over the significantly cheaper TF-IDF+SVM baseline (0.937 vs. 0.928 macro F1) statistically significant. The SVM trains and predicts orders of magnitudes faster (see Table 8) at a fraction of the storage cost. This is an important and currently under-emphasized finding for a paper that proposes AAAC as a practical moderation benchmark. It suggests that a strong lexical baseline is competitive with fine-tuned transformers for this dataset size and task. Classes with the least training support (308 and 368 samples respectively for tool_misuse and sensitive_request) are the weakest ones, and are the most semantically close to safe (a sensitive-sounding query could be a legitimate request or a privacy-violating one depending on framing rather than vocabulary). This should explain the lower precision/recall observed.

Key Findings

For 5-way intent classification, DziriBERT achieves the highest macro F1 (0.937-0.940), but it has a statistically insignificant advantage over a TF-IDF+SVM baseline. The two smallest and most context dependent classes, tool_misuse and sensitive_request, perform the worst.

5.2 RQ2: Binary Adversarial Detection

Experimental Setup

The same five models were trained on the binary is_adversarial label, using identical data splits and preprocessing.

Observations

Table 5 shows that all five models achieve substantially higher macro F1 on

the binary task than on the 5-way task. DziriBERT again ranks first (0.987), followed by MARBERT (0.968), SVM (0.948), AraBERT (0.945), and Random Forest (0.938). McNemar's test between DziriBERT and MARBERT (the top two) gives b = 9, c = 3, p = 0.1460 – not significant.

Table 5: Binary adversarial-detection leaderboard, test set (n = 309).

Model	Acc.	Prec.	Rec.	Macro F1
DziriBERT	0.987	0.987	0.987	0.987
MARBERT	0.968	0.969	0.968	0.968
SVM	0.948	0.948	0.948	0.948
AraBERT	0.945	0.946	0.945	0.945
RandomForest	0.939	0.943	0.938	0.938

DziriBERT's per-class results are near-symmetric and near-ceiling (benign: P=0.987, R=0.987, F1=0.987, n = 155; harmful: P=0.987, R=0.987, F1=0.987, n = 154), indicating no systematic bias toward either class.

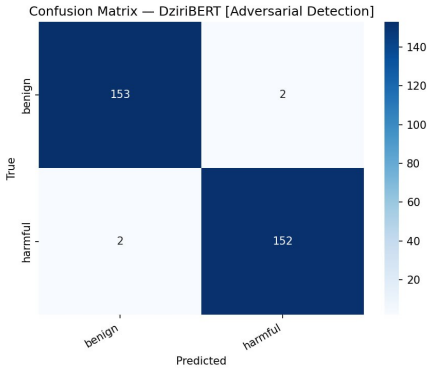


Fig. 6: Confusion matrix of DziriBERT on the binary adversarial-detection task, test set (n = 309)

Discussion

Combining the four adversarial categories into one label makes the task considerably simpler. Each model's macro F1 improves relative to its 5-way outcome (e.g. MARBERT: 0.908 → 0.968; AraBERT: 0.909 → 0.945). This is expected, as binary detection requires only the distinction between safe and any other non-safe pattern, not the distinction between four adversarial sub-types that may share a surface vocabulary (e.g. sensitive_request and hate_speech both sometimes target a named individual). This has a direct practical implication for deployment, with a lightweight binary gate-keeping classifier perhaps sufficient

for coarse moderation, with finer grained categorization left for flagged content. As with RQ1, the difference in performance between DziriBERT and the runner-up is not statistically significant, suggesting that dialect-specific pretraining is helpful but not conclusively decisive at this test-set size (n = 309).

Key Findings

For all models evaluated binary adversarial detection is significantly easier than five-way intent classification (macro F1 increases of 0.03-0.06). DziriBERT (macro F1 0.987) is still the top performer, but it is no longer significantly ahead of MARBERT.

5.3 RQ3: LLM Zero-Shot / Few-Shot Baseline vs. Fine-Tuned Encoders

Experimental Setup

The Llama-3.1-8B-Instant (Groq API) was tested on the same 309-row test set with a system prompt describing the AAAC label taxonomy. The test set had zero in-context examples (zero-shot) and 2 examples/class/script (intent task) or 3 examples/class (adversarial task) (few-shot). We did not perform any gradient-based fine-tuning.

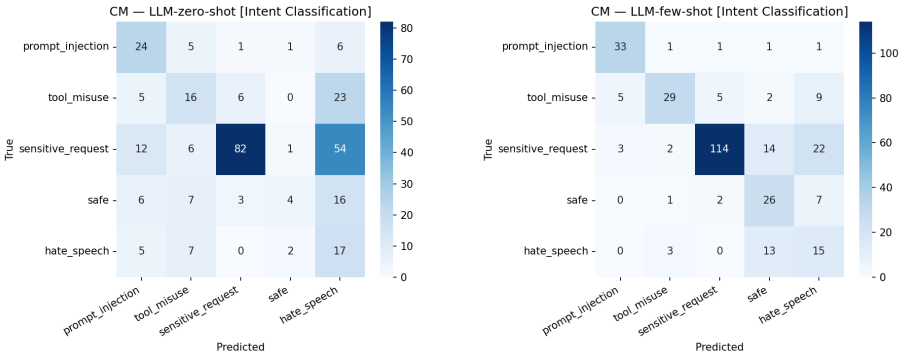
Observations

Table 6 situates the LLM baseline against the fine-tuned-encoder leaderboard. On both tasks, the LLM substantially trails every fine-tuned model, even the classical SVM baseline; few-shot prompting yields a large improvement over zero-shot on both tasks (intent: +0.259 macro F1; adversarial: +0.204 macro F1), but does not close the gap to fine-tuned encoders.

Table 6: LLM baseline vs. best fine-tuned model, by task.

Model	Intent Macro F1	Adv. Macro F1
DziriBERT (best fine-tuned)	0.937	0.987
SVM (best non-transformer)	0.928	0.948
LLM few-shot	0.652	0.793
LLM zero-shot	0.394	0.589

For the intent task, the zero-shot LLM has very low precision/recall on tool_misuse (F1 = 0.23) and sensitive_request (F1 = 0.18), while performing comparatively better on safe (F1 = 0.66, though recall is only 0.53, i.e. it over-predicts adversarial labels). Few-shot prompting substantially improves hate_speech (F1 0.54 $\rightarrow$ 0.85) and sensitive_request (F1 0.18 $\rightarrow$ 0.57), but tool_misuse remains the weakest class even with few-shot examples (F1 = 0.35). For the adversarial task, few-shot prompting yields a highly asymmetric result: precision on the benign class reaches 0.97 but recall is only 0.62, while the harmful class shows the inverse pattern (precision 0.72, recall 0.98) – the few-shot LLM is biased toward predicting "harmful."



(a) Zero-shot (b) Few-shot

Fig. 7: Confusion matrices for the LLM baseline (Llama-3.1-8B-Instant) on the five-way intent classification task, test set (n=309); (a) zero-shot, (b) few-shot.

Discussion

The large gap between the LLM baseline and fine-tuned encoders implies that AAAC should be considered primarily as a fine-tuning benchmark, and not as a prompting benchmark. For example, a general-purpose 8B model not trained on this task performs much worse as a classifier on this dialect and label taxonomy than a 110M-parameter fine-tuned encoder. The asymmetric precision/recall pattern for the few-shot adversarial result (high precision, low recall for "benign") indicates that the LLM is conservatively biased towards labeling Algerian Arabic content adversarial when it is unsure. This is a reasonable behavior because the LLM's pretraining has limited Algerian Arabic data, increasing the probability that benign-but-unfamiliar phrases are misclassified as suspicious. This pattern is a safety-relevant discovery in and of itself: a general-purpose LLM that is miscalibrated to overflag legitimate requests in Algerian Arabic and deployed as a moderation layer has real usability costs for Algerian Arabic speakers.

This experiment tests the LLM as a classifier by asking it to output a label, rather than testing the LLM's behavior in refusing adversarial prompts in a chat-like interaction for end-users. We state a significant limitation in scope for this RQ and for the paper's assertions as a whole. These evaluation settings are related but are distinct paradigms. So the contributions of this paper are to be understood in the context of LLM-based classification and not in the context of conversational safety or refusal behavior.

Key Findings

A general-purpose 8B LLM underperforms fine-tuned Arabic encoders significantly for both AAAC tasks even with few-shot prompting (intent: 0.652 vs. 0.937 macro F1; adversarial: 0.793 vs. 0.987). Few-shot adversarial detection result is consistent with a conservative bias to over-flagging adversarial Algerian Arabic content, showing the precision/recall asymmetry.

5.4 RQ4: Per-Script Performance Breakdown

Experimental Setup
For each of the five fine-tuned/classical models and both tasks, we partitioned the saved test-set predictions by language_variant and recomputed macro F1 on each subset (Algerian Derja: n = 102; Arabizi: n = 126; Franco-Arabic: n = 81), without retraining.
Observations
Table 7 shows a consistent pattern across nearly every model on the intent task: Arabizi is the hardest script, with macro F1 dropping 6–10 points relative to Algerian Derja for every transformer model (MARBERT: 0.963 → 0.810; AraBERT: 0.968 → 0.837; DziriBERT: 0.960 → 0.893). The SVM baseline shows the smallest Arabizi gap (0.930 → 0.902) of any model on the intent task. The same Arabizi-is-hardest pattern holds on the adversarial task, though the absolute gaps are much smaller given the higher overall ceiling (e.g. DziriBERT: 1.000 algerian_derja vs. 0.976 arabizi).

Table 7: Per-script macro F1, intent task, test set.

Model	Derja	Arabizi	Franco-Ar.
MARBERT	0.963	0.810	0.942
AraBERT	0.968	0.837	0.898
DziriBERT	0.960	0.893	0.942
SVM	0.930	0.902	0.929
RandomForest	0.875	0.789	0.885

Discussion
The Arabizi performance gap is the most prominent script-dependent robustness discovery in this paper and is directly relevant to AAAC's safety motivation. The Latin-script, digit-augmented orthography of Arabizi (e.g., 3, 7, 9 for Arabic phonemes) is underrepresented in the pretraining corpora of all three transformer encoders, none of which were pretrained with substantial Arabizi exposure. While the Derja-vs-Arabizi gap is the smallest for the transformers (6.7 points), this shows that its pretraining corpus is still predominantly Arabic-script even though it is dialect-matched. This is consistent with the observation that TF-IDF over character/word n-grams is less sensitive to script identity than a subword tokenizer trained primarily on Arabic-script vocabulary, as seen in the comparatively flat performance of the SVM across scripts. This finding has an immediate, actionable implication: a moderation system based on standard Arabic-pretrained encoders is quantifiably easier to circumvent by writing adversarial content in Arabizi than in Arabic script. This is precisely the sort of script-level robustness gap that AAAC aims to reveal.
Key Findings
Across most models and both tasks, Arabizi is consistently the hardest script, with macro F1 drops of 6-10 points against Derja in Algerian-script on the intent

task. This is a script-dependent robustness gap with direct safety implications: These classifiers struggle more to detect adversarial content written in Arabizi.

5.5 Efficiency Comparison

Table 8 provides a comprehensive overview of the training time and model size for all ten model-by-task configurations. The transformer models take 130-250 seconds and 500-650 MB respectively, while the TF-IDF+SVM baseline is trained in less than 0.05 seconds and less than 1 MB. This shows a large efficiency-accuracy trade-off as SVM's macro F1 is only 0.9-3.9 points behind the best transformer depending on the task.

Table 8: Training time and model size by model and task.

Task	Model	Time (s)	Size (MB)
Intent	MARBERT	141.4	651.4
Intent	AraBERT	195.3	540.8
Intent	DziriBERT	143.8	497.8
Intent	SVM	0.04	0.17
Intent	RandomForest	4.71	43.5
Adversarial	MARBERT	249.9	651.4
Adversarial	AraBERT	207.3	540.8
Adversarial	DziriBERT	192.9	497.8
Adversarial	SVM	0.03	0.03
Adversarial	RandomForest	3.78	26.7

6 Discussion

The four RQs collectively support three main claims. Dialect-matched pretraining (DziriBERT) yields the best central-estimate performance on both AAAC tasks. But at this test-set size, this advantage is not statistically significant relative to a general-dialect transformer (MARBERT, adversarial task) or a lexical baseline (SVM, intent task). Therefore, any strong claim that dialect-specific pretraining is necessary to strong performance on AAAC is tempered. Second, AAAC draws a clear line between fine-tuning and prompting approaches. A competitive general-purpose LLM that is used only as a zero/few-shot classifier lags behind fine-tuned encoders by 14–34 macro-F1 points, depending on the task. This suggests that dialect-specific content moderation in Algerian Arabic currently benefits substantially from task-specific fine-tuning rather than the in-context capabilities of a general-purpose model.

Third, the per-script analysis (RQ4) shows that aggregate metrics hide a real robustness gap: every model studied is measurably weaker on Arabizi than on Arabic-script Derja. This is the very kind of script-level blind spot that an

adversarial actor could exploit, and that a benchmark limited to a single script would never reveal. We believe this is the most unique empirical contribution of AAAC towards single-script Arabic NLP benchmarks.

These results, considering LLM safety, indicate that the use of a general-purpose LLM's own judgment (zero/few-shot classification) for moderation of Algerian Arabic content is not reliable at this time. Additionally, the precision/recall asymmetry from the few-shot adversarial result (Section 5, RQ3) presents a usability concern: an LLM that is unsure about unknown Algerian Arabic phrasing appears to lean towards classifying it as suspicious, which could disproportionately impact legitimate Algerian Arabic speakers of moderation-gated systems.

7 Limitations

- *Single annotator.* One annotator performed all manual checks and corrections. Labeling decisions on inherently ambiguous categories (especially the safe/sensitive_request and hate_speech/sensitive_request boundaries) are based on the judgment of one person, not a consensus process, and inter-annotator agreement could not be computed. Before using AAAC as a shared benchmark in the form of a leaderboard, we recommend a multi-annotator validation pass with reported κ.
- *Small per-class test support.* In the test set, per-class support is limited to 309 rows, with values between 31 (tool_misuse) and 155 (safe). The per-class precision/recall at this scale is subject to a great deal of sampling variance. Some of the confidence intervals reported in Section 5 are 6-7 points wide. Note that the results on the smallest classes (sensitive_request, tool_misuse) should be considered indicative of the trend but not exact.
- *Scope: classifier benchmark, not live refusal-behavior evaluation.* The results provided in this paper are solely related to supervised text classifiers and an LLM used as a zero/few-shot classifier. They do not measure whether a deployed conversational LLM would actually obey or refuse the adversarial prompts in AAAC when exposed to them in an actual chat interaction. A natural and valuable follow-up is to evaluate the actual model compliance/refusal rates on these prompts, but outside the scope of this paper. AAAC's prompt_injection and tool_misuse categories were specifically designed with adversarial-prompting research in mind.
- *Coverage of Algerian Arabic variation.* AAAC includes three written-script varieties, but does not consider regional linguistic variation within Algeria (e.g., dialectal variation, such as Wahrani, Chaoui and other local dialects, or Amazigh language varieties), register (formal/informal), or speaker demographics, all of which may impact generalization.
- *Single LLM baseline.* We evaluated only one general-purpose LLM (Llama-3.1-8B-Instant) as a no-fine-tuning baseline. These results may not generalize to larger models or differently trained models.

8 Future Work

- Multi-annotator re-validation of a representative subsample with reported κ to improve quality claims for shared-benchmark use.
- Live evaluation of the refusal/compliance behavior of deployed conversational LLMs on AAAC's adversarial prompts, which complements the classifier-based evaluation in this paper.
- Data collection designed to address the performance gap found in RQ4, such as Arabizi-centric pretraining data or script normalization preprocessing before classification.
- Evaluation of other general purpose LLMs (different sizes, different vendors) as zero/few-shot baselines.
- The two categories with the lowest per-class F1 and the smallest support in the current release, tool_misuse and sensitive_request, will be extended.

9 Conclusion

We introduce AAAC, a dataset of 3,090 samples in Algerian Arabic, with five safety-relevant intent classes and three script variants, and evaluate five classifiers and a general-purpose LLM baseline on intent classification and binary adversarial detection. The best results on both tasks are obtained by a dialect-specific pre-trained encoder (DziriBERT) with macro F1 of 0.937 and 0.987 respectively, but it is not significantly better than strong alternatives at this test-set size.

A general-purpose LLM used as a zero/few-shot classifier underperformed significantly compared to all fine-tuned models. A per-script analysis reveals that all models studied are measurably weaker on Arabizi than on Arabic-script Algerian Derja, exposing a robustness gap of direct relevance to AAAC's motivating safety use case. We have released the AAAC, its fixed train/validation/test splits, and the accompanying evaluation code to facilitate reproducible follow-up work on Algerian Arabic safety and moderation classifiers. available at https://github.com/AAAC-Corpus/AAAC

References

1. zachz.: prompt-injection-benchmark [Dataset]. HuggingFace. https://huggingface.co/datasets/zachz/prompt-injection-benchmark
2. Antoun, W., Baly, F., Hajj, H.: AraBERT: Transformer-based Model for Arabic Language Understanding. https://arxiv.org/pdf/2003.00104
3. Abdul-Mageed, M., Elmadany, A., Nagoudi, E.M.B.: ARBERT & MARBERT: Deep Bidirectional Transformers for Arabic. https://arxiv.org/pdf/2101.01785
4. Abdaoui, A., Berrimi, M., Oussalah, M., Moussaoui, A.: DziriBERT: a Pre-trained Language Model for the Algerian Dialect. https://arxiv.org/pdf/2109.12346
5. Bender, E.M., Friedman, B.: Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science. TACL, 2018. https://aclanthology.org/Q18-1041.pdf

6. Darwish, K., et al.: A Panoramic Survey of Natural Language Processing in the Arab World. Communications of the ACM, 64(4), 72–81 (2021). https://arxiv.org/pdf/2011.12631
7. Osact5. https://aclanthology.org/2022.osact-1.20.pdf
8. Ganguli, D., et al.: Red Teaming Language Models to Reduce Harms: Methods, Scaling Behaviors, and Lessons Learned. https://arxiv.org/pdf/2209.07858